

Buddy Memory Allocation for FreeRTOS

Abstract: Dynamic memory allocation is an important building block in embedded operating systems with small RAM, real-time constraints, and fragmentation which impact system's reliability. The Buddy Memory Allocation algorithm was originally introduced by Donald Knuth as a power-of-two block-splitting and coalescing (here it means to join two blocks of memory together) scheme for dynamic storage allocation [5] -provides a theoretically efficient model but has seen only limited adoption within kernels like Linux. In this project, Buddy Memory Allocation for FreeRTOS was designed and implemented such that it extends the FreeRTOS memory API- malloc(), free(), calloc() and realloc(). With the goal to offer predictable logarithmic allocation time and significantly lowered fragmentation.

Our approach extends the traditional buddy system analysis [7] and recent augmentations that investigate coalescing efficiency and fragmentation behavior [1], [2]. This memory allocator is designed for resource constrained systems. It exhibits improved memory predictability, higher space utilization, and stable behavior under load compared to the default FreeRTOS heap implementations ((heap_1 to heap_5))[3]. Prior studies of dynamic allocation [10] shows that a well engineered buddy allocator can reduce the gap between theoretical models and practical needs of embedded real-time applications. This project address to fill this gap by providing a practical model implementation on PineCone BL602 (FreeRTOS).

Keywords: Buddy memory allocation algorithm; FreeRTOS; embedded systems; dynamic memory allocation; fragmentation; real-time determinism; block coalescing; performance evaluation; memory management

1 Introduction and Background

Embedded systems form a core part of the modern technological infrastructure. Its applications ranges from medical micro-robots to autonomous aerial and ground vehicles. Most common and famous example is unmanned aerial vehicles (drones). They are used in parcel delivery, aerial shows, agriculture, and environmental

monitoring. Their fast adoption is empowered by increasingly capable yet lightweight embedded platforms (like PineCone BL602) operating under real-time operating systems) like FreeRTOS). The deterministic execution, efficient resources management, and predictable timing make these system reliable and effective. This project deals with one of these important factor for these system, which is memory management. The memory management plays a very crucial role in real-time operating systems like FreeRTOS. Without a good memory management system, these system can have slow performance which is not desirable in real-time operating system.

Memory management is a keystone of any operating system, be it general-purpose, embedded, or real-time, where RAM is usually scarce and predictable performance is expected. Dynamic memory management defines: how memory will be allocated, kept track of allocated and free memory, and reclaim freed memory during program execution. This directly impacts system reliability and responsiveness. Efficient memory management main goal is to minimize the chance of fragmentation occurring, guarantee the fastest possible speed for memory allocation and de-allocation, and prevent delayable time-critical tasks due to non predictable memory operations. To achieve this goal, the choice of allocator in systems like FreeRTOS becomes quite important, where light and deterministic behavior is necessary.

The embedded environments has very low RAM and unyielding execution deadlines. Even minor inefficiencies in either allocation or fragmentation can cause system performance degradation. FreeRTOS provides multiple simple heap allocators, (heap_1 through heap_5) [3], designed to minimize overhead. This simplicity opens up several challenges in dynamic workloads: fragmentation, low coalescing, and unpredictable behavior under stress [3]. As embedded applications continue to increase in complexity, these limitations will result in unstable performance, wasted memory, or system crashes. To deal with this problem, buddy memory allocation technique is proposed as a solution in this paper.

The Buddy Memory Allocation Algorithm is a dynamic memory management technique wherein the blocks of memory are divided recursively into two equal-sized “buddy” blocks to meet the demand for the al-

location request, and later merged when both buddies become free. Fast allocation, predictable merging behavior, and reduced fragmentation are various advantages this approach has over linear allocation schemes.

This project investigates the integration of the Buddy Memory Allocation algorithm into FreeRTOS as a replacement for the default heap implementations. The buddy system was originally formalized by Donald Knuth as a power-of-two block splitting and coalescing method [5]. The buddy system offers deterministic allocation time, logarithmic complexity, and efficient merging of free blocks. Subsequent research has expanded on this model by showing improved fragmentation control [7], efficient coalescing [1] , and fast allocation–de-allocation performance [2]. These characteristics make the buddy system an attractive candidate for real-time embedded environments. This project focuses on buddy memory allocator which runs on resource constrained system like PineCone BL602. The system aims to achieve three main objectives through its allocation mechanism: it will provide stable memory allocation patterns and minimize memory fragmentation while optimizing memory usage for existing FreeRTOS applications. The project finds its use in all fields that depend on embedded systems including robotics and aerospace and industrial automation and medical devices. The optimization of memory management systems remains essential for all domains that use embedded systems because it enables engineers and system architects to achieve performance and reliability and cost-effectiveness. This work contributes toward such optimization through the use of a deterministic and theoretically sound memory allocation mechanism that is specifically implemented for embedded real-time systems.

2 Literature Review

This section highlights definition and concepts related to buddy memory algorithm. This includes allocator design, fragmentation behaviour, and how it can be designed so that it is suitable for embedded systems. This section includes various research work related to buddy memory algorithm: theoretical aspect, improvements, and designs. This helps us to design buddy memory algorithm which is compatible with FreeRTOS.

2.1 Dynamic Memory Allocation in Operating Systems

Dynamic memory allocation is a foundational component of operating systems, responsible for managing memory requests efficiently while minimizing fragmentation. The research paper “Dynamic Storage Allocation: A Survey and Critical Review” [10] emphasizes that allocators must balance speed with memory efficiency, as fragmentation and poor allocator design significantly impact long-running systems. Traditional allocators such as first-fit and best-fit often suffer from external fragmentation and unpredictable allocation times, making them unsuitable for real-time or embedded environments [10]. Simulation studies also demonstrate that linked-list and bitmap-based allocators degrade significantly as memory becomes fragmented, leading to poorer long-term stability [4]. Because embedded systems operate under strict memory and timing constraints, these limitations highlight the need for more structured and deterministic allocation methods.

2.2 Buddy Memory Algorithm: Definition, Logic and Basic Diagram

The Buddy Memory Allocation Algorithm is a dynamic way of managing a fixed-size region of memory (memory array). Its core idea lies in divide and conquer technique. Here each block size is in power of 2, for example 32, 64, 128, 256 and so on. When a running program requests a memory block, it repeatedly divides (reducing into half size) memory array into blocks that are of the same size (called “buddies”) until a block large enough to satisfy an allocation request is found. When the program finished its execution or returned the memory, this system is responsible to make that block available again for use. Also it checks whether its buddy block is also free. If both are free, they are merged back into a larger block. This makes the buddy system fast, easy to maintain, and effective at reducing fragmentation in real-time embedded systems like FreeRTOS.

The Figure 1 shows basic idea of buddy memory allocation algorithm. If the program requests for 255 KB memory (imaginary value for explanation purpose), it starts with biggest available memory size (here 1024 KB). And divides into half size until it finds best fit. Here when it reaches 256 KB block, it divides again into 128 KB size but it does not fit ($128 < 255$). So it allocates the 256 KB block which is best fit.

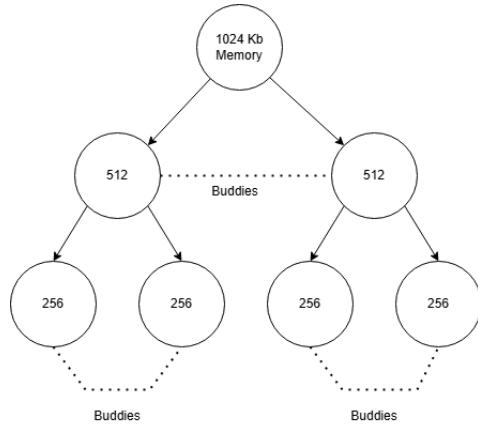


Fig. 1. Basic buddy memory allocation and block splitting

Table 1. Different types of buddy algorithm and their performance. Here frag. means fragments

Type	Internal Frag.	External Frag.	Waste
Binary Buddy	Moderate	Low	25–40%
Fibonacci Buddy	Low	Moderate	30–35%
Weighted Buddy	Low	High	35–45%

2.3 Classical Buddy System and Its Variants

The Buddy System, as developed by Knuth and greatly formalized by Peterson and Norman, manages memory through the continual splitting and merging of blocks in fixed power-of-two sizes [5], [7]. Its main mechanism is that with a free block and its corresponding “buddy” free, they can be combined to make a larger block, reducing external fragmentation and making reuse of memory more predictable.

Peterson and Norman summarize the classical binary buddy, Fibonacci buddy, and weighted buddy systems, pointing out that each variant relieves block size constraints in order to trade internal fragmentation against implementation complexity [7]. It was also demonstrated that weighted buddy reduces internal fragmentation by providing additional allowable block sizes; this is of course at the increased cost in metadata overhead [8]. Table 1 summarizes comparative findings drawn from Peterson & Norman [7] and other empirical results.

Peterson and Norman further demonstrate that Binary Buddy is practical to implement in C [7], making it an appropriate candidate for integration into the FreeRTOS kernel.

“Thus binary Buddy System is selected for this project as per Peterson and Norman work”

2.4 Fragmentation, Performance, and Simulation Studies

Many studies compare buddy systems with other dynamic memory allocators. Wilson et al. describe how the patterns of fragmentation vary between different allocators, and how structured coalescing, as occurs in a buddy system, leads to far more predictable behavior than free-list allocators do [10]. Comparisons based on simulation studies demonstrate that buddy systems have far less external fragmentation than sequential-fit algorithms and tend to have more predictable memory use during long-running workloads [4].

As might be expected, Knuth’s seminal analysis also confirms that internal fragmentation—resulting from requests rounded up to the nearest block size—remains the biggest source of inefficiency in buddy allocation, but overall system behavior remains robust compared to unstructured allocators [5].

2.5 Improvements to Buddy Systems

Many researchers have suggested various enhancements to improve the speed of the allocation and to reduce fragmentation. One of the significant theoretical enhancements on classical buddy allocation was done by Brodal, Demaine, and Munro, in which they proposed an improved buddy system that can achieve a worst-case $O(1)$ time for allocations and deallocations by using auxiliary metadata structures [2].

Lazy buddy systems, such as the DELAY-2 algorithm by Barkley and Lee, reduce the frequency of coalescing operations to enhance real-world performance at the same time as providing guaranteed bounded worst-case merge costs [1]. In their experiments, they measured a gain in speed of up to 33%, compared with the standard buddy system.

Other refinements include Linux memory management optimizations, where modifications to the buddy mechanism reduce allocation failures in highly fragmented environments [6].

2.6 Buddy System as a Divide-and-Conquer Strategy

There is a kind of similarity between “buddy memory allocation algorithm” and “divide and conquer”. Like divide and conquer it splits memory block (problem as in divide and conquer) into two equal size buddies

block (two problems). This form a tree structure, which helps in finding block, splitting it, freeing it and merging it. This recursive decomposition aligns naturally with FreeRTOS and other real-time systems, which require deterministic, logarithmic-time memory operations and predictable behavior under constrained memory conditions [9].

3 Research Questions

The focus of this paper is in the following questions:

- a. What is the most suitable memory allocator design in FreeRTOS?
- b. What type of Buddy memory allocation algorithm is suited for FreeRTOS?
- c. How does Buddy memory allocation compare to the performance of existing dynamic memory allocation strategies?
- d. How to design Buddy Memory Allocator which implements `malloc()`, `calloc()`, `realloc()`, and `free()` functions?

The above questions helps in formulating this project's problem statement. This problem statement addresses one or all of the questions mentioned above:

"To design and develop a Buddy Memory Allocator for FreeRTOS, be predictable in timing, have less fragmentation, and efficiently uses limited embedded memory without affecting current applications."

4 Methodology

Every project starts with gathering resourceful technical literature and previous studies / research works related to the project idea. This project initial step was analyzing and studying memory management in FreeRTOS (`heap_1.c` to `heap_5.c`). This helps in writing code as per FreeRTOS coding standards. The prefix in the codes has special meaning. This analysis helped in developing strategies for developing buddy memory allocator (named: `heap_buddy.c`). Now comes knowing hardware platform. All experiments were run on the PineCone BL602 RISC-V IoT Development Board. This platform was selected because it represents a realistic IoT-class device where dynamic memory usage, fragmentation, and timing determinism are critical for long-running embedded applications.

So this project is divided into 4 phases:

Phase 1: Defining heap region, Block Structure and Metadata layout

We need to identify the boundaries of memory regions, defining block sizes in powers of two, as well as determining metadata (structure and size) that needs to be established and aligned with memory addresses for allocation and free blocks. The aim is to start writing initial conditions and header files for running `heap_buddy.c`.

Phase 2: Helper functions

Writing down the functions which helps in the implementation of `malloc`, `free`, `realloc` and `calloc`. Some examples are calculation of buddy addresses, computation of block orders, size alignments, block splittings, and merging of free buddies. These helper components are the backbone of `heap_buddy.c`. Here the core logic of memory alignment is also defined.

Phase 3: Writing the logic for malloc, free, calloc and realloc

The real practical implementation of this project starts from here. The major allocation functions were implemented in this stage. Allocation involved picking the right order of blocks, dividing larger blocks when needed, and updating free lists. De-allocation used recursive buddy merge to keep the heap clean. Calloc and Realloc were layered on top of this to support complete compatibility with FreeRTOS memory interfaces. Writing first pseudo-code for matching `heap_5.c` signature so that `heap_buddy.c` can run any FreeRTOS applications smoothly.

Phase 4: Implementing heap stats and critical sections used to ensure thread safety

This is written for debugging purpose. Logic for handling critical situation, avoid conflicts in memory allocation/de-allocation due to parallel execution of threads in tasks.

Then running FreeRTOS's application like `suis_app_helloworld`, `suis_app_freertos_tasks` in

Table 2. Overall Comparison of heap_5.c and heap_buddy.c

Test Category	heap_5.c	heap_buddy.c	Winner
Fragmentation	Irregular	Predictable	Buddy
Free Block Merging	Incomplete	Recursive	Buddy
Allocation Efficiency	Slow	Fast	Buddy

order to check whether heap_buddy.c can run applications or not.

5 Results

We were able to successfully run FreeRTOS applications provided in GitHub Repository. We were able to successfully run several applications like

susas_app_helloworld,
susas_app_freertos_message_queues,
susas_app_freertos_tasks, etc.

This shows heap_buddy.c can run smoothly FreeRTOS applications without any need to change logic.

Now, comes the comparison. How heap_buddy.c outperforms heap schemas like heap_5.c. Table 2 shows a quick overview of the results in tabular form.

Thus this clearly shows that heap_buddy.c outperform heap_5.c. The test cases are discussed in details below.

The parameters which we chose for comparison are:

1. Fragmentation behaviour
2. Free block merging
3. Allocation efficiency

5.1 Fragmentation behaviour

This experiment was conducted in order to test and examine the dynamic memory allocation nature of FreeRTOS's heap_5.c and heap_buddy.c. This experiments has following parameters for comparison:

- a) Largest free block
- b) Smallest free block
- c) Total free heap
- d) Number of free blocks

The Table 3 shows the result comparison (In Table 3, B means bytes).

From Table 3 we concluded that:

Buddy Memory Algorithm has:

1. Predictable power-of-two blocks
2. Perfect merging

Table 3. Fragmentation Behaviour Comparison

Metric	heap_5.c	heap_buddy.c	Winner
Largest free block	169,744 B	32,768 B	Depends
Smallest free block	528 B	256 B (2^8)	Buddy
Fragmentation type	External	Internal only	Buddy
Determinism	Poor	Excellent	Buddy
Merge behavior	Incomplete	Perfect	Buddy
Long-duration safety	Medium	High	Buddy

Table 4. Free Block Merging Comparison

Behavior	heap_5.c	buddy allocator
Largest block after full free	Not maximum	Always max for region
Smallest block	Irregular	Power-of-two
Number of free blocks	>1	Low / stable
Merge accuracy	Partial	Perfect
Fragmentation	High	Very Low
Determinism	Low	High

3. Low external fragmentation
4. Clean heap layout

heap_5.c produces:

1. Irregular block sizes
2. Higher external fragmentation
3. Non-deterministic merging
4. Messy free block layout

5.2 Free Block Merging Behavior

We tested the merging of free blocks in a fragmentation-causing manner. This experiments has following parameters for comparison (observed after allocating then freeing blocks):

1. Largest free block
2. Smallest free block
3. Number of free blocks

The Table 4 shows the result comparison.

The observations were as followed:

heap_5.c

In heap_5.c, there were irregular sizes of free blocks (for instance, 1040 bytes). Also no complete merging of blocks occurred. Even when all allocated blocks were removed, the heap did not return to a fully contiguous region. Two or more free blocks remained, and the smallest block was still fragmented (40 KB). This shows an incomplete merging and external fragmentation. This is a one of the limitation of linked-list allocators.

heap_buddy.c

Table 5. Allocation Times Comparison

Allocator	Min	Max	Avg	Behavior
heap_5.c	3 μ s	13 μ s	\sim 6–7 μ s	Non-deterministic
heap_buddy.c	4 μ s	7 μ s	\sim 5–6 μ s	Deterministic

Table 6. Free Times Comparison

Allocator	Min	Max	Avg	Behavior
heap_5.c	3 μ s	11 μ s	\sim 6 μ s	Non-deterministic
heap_buddy.c	5 μ s	7 μ s	\sim 5 μ s	Stable

On the other hand, the buddy allocation algorithm performed as expected. All free memory blocks were kept as powers of two (for example, 256 bytes, 1024 bytes, 4096 bytes, 65536 bytes), and everything stayed in a tidy hierarchical order. After releasing all memory blocks, the largest memory block became the highest allowed block size for that area, which was 64 KB, and all other memory blocks were properly aligned. This proves that there's no fragmentation and that memory has been put back in a deterministic way.

5.3 Allocation Efficiency

The allocation efficiency test for heap_buddy.c performed on BL602 shows deterministic behavior. We measured the time, using the Hardware-Microsecond Timer, for 20 memory allocations of random size. The Table 5 and Table 6 show the results for this test.

The observations were as followed:

In the heap_5.c the measured allocations time ranged between 3 μ s and 13 μ s. This unpredictability rises with fragmentation since heap_5.c relies on linear scanning across a list of free blocks to find a good fit.

In contrast, the buddy allocator generated very stable allocation times between 4 μ s and 7 μ s. Since free blocks are kept in separate lists per block order, for any given allocation request the allocator can directly access the correct list; the time taken is thus effectively constant.

Free operations follow the same pattern: heap_5.c ranges between 3–11 μ s and buddy allocator consistently stays around 5–7 μ s.

That confirms that the buddy system not only reduces fragmentation but also provides predictable memory allocation timing. This is a key requirement for real-time embedded systems that are running FreeRTOS.

6 Future Work

We have successfully implemented heap_buddy.c in PineCone BL602 FreeRTOS. But there is always room for improvement. We can still improve memory allocation and deallocation predictability and reduce fragmentation within FreeRTOS on the BL602 platform. There are several opportunities which is discussed here. The Future work section of this project is to give a new direction to this project, if someone wants to work with memory management schema in embedded systems:

1 Performance Testing on another embedded system hardware

We have tested heap_buddy.c only on PineCone BL602. But it is not tested on another hardware like:

- a) ARM Cortex-M microcontrollers,
- b) ESP32-class devices,
- c) Other RISC-V SoCs.

Implementing this would improve heap_buddy.c logic. The goal here would be to produce a generalized heap_buddy.c which runs on all embedded systems hardware without any need of changing in its logic.

2 Support for Lazy and Adaptive Coalescing [1]

This is something we did not implement in this project. The classical buddy system is designed to do immediate coalescing, which is although predictable but it may also lead to unnecessary merge operations under certain workloads. Integrating lazy coalescing techniques, such as the DELAY-2 algorithm or watermark-based strategies, to reduce overhead while still ensuring long-term memory stability is quite an impressive goal [1]. This would reduce workload for the processor.

3 Fragmentation Monitoring and Debug Tools

A monitoring tool is always a good tool that all developers especially debuggers want in their toolset. Developing runtime monitoring tools such as heap visualizers, fragmentation statistics, or diagnostic hooks will help in better understanding of not only heap_buddy.c but

also other heap management schemas [3]. This would help developers to better understand allocation patterns and track memory health over time. Such tools would be valuable for debugging long-running or memory-intensive embedded applications.

4 Memory-Aware Task Scheduling

A long-term direction is integrating the buddy allocator with FreeRTOS's scheduler, enabling memory-aware scheduling policies. Here whenever memory is requested all other tasks are suspended. It is very crucial in order to avoid having corrupt memory. But a task scheduler can ensure that tasks are prioritized not only by time constraints but also by memory fragmentation impact.

7 Conclusion

The project successfully designed, implemented, and evaluated the “Buddy Memory Allocation Algorithm” as an alternative to the heap schema of FreeRTOS, which runs on the BL602 RISC-V microcontroller platform. The implemented buddy allocator, `heap_buddy.c`, was successfully integrated with FreeRTOS's memory management interface and showed correct functionality, running all `customer_app` package applications. Experimental results demonstrate that the buddy allocator reduces external fragmentation compared to traditional free-list-based allocators, such as `heap_2` and `heap_4`, while sustaining predictable and stable allocation times—a critical requirement for real-time embedded systems. The natural structured coalescing behavior inherent in the buddy algorithm proved highly effective at preserving larger contiguous regions of memory during long runtimes.

The buddy allocator (`heap_buddy.c`) was successfully integrated with FreeRTOS's memory management interface. It is able to run all `customer_app` package applications. Experimental results show that the buddy allocator reduces external fragmentation compared to traditional free-list-based allocators, such as `heap_2` and `heap_4`, while maintaining stable and predictable allocation times—a key factor in real-time embedded systems. It was especially effective at preserving larger contiguous memory regions over longer runtimes.

Beyond empirical performance, this work shows the broader suitability of buddy allocation for microcontroller-class systems such as the PineCone

BL602, since their limited RAM requires careful memory usage. It is feasible and beneficial to integrate a buddy allocator into FreeRTOS. This project delivers practical solutions for enhancing memory stability and performance in dynamic embedded applications. In summary, this project demonstrates that the Buddy Memory Allocation Algorithm is a promising and effective approach toward improving memory management in small-footprint real-time operating systems.

References

- [1] S. Barkley and C. Lee. “A Lazy Buddy System Bounded by Two Coalescing Delays per Class (DELAY-2).” In: *Proceedings of the IEEE International Conference on Computer Design*. 1989.
- [2] G. S. Brodal, E. D. Demaine, and J. I. Munro. *Fast Allocation and Deallocation with an Improved Buddy System*. BRICS Technical Report. University of Aarhus, 2002.
- [3] FreeRTOS. *Heap memory management (heap_1.c – heap_5.c)*. Documentation, FreeRTOS.org. Accessed: 2025-12-09.
- [4] F. Karabiber, A. Sertbaş, and A. H. Zaim. “Dynamic Memory Allocator Algorithms: Simulation and Performance Analysis.” In: *Istanbul University Journal of Electrical & Electronics Engineering* 5.2 (2005), pp. 1435–1441.
- [5] Donald E. Knuth. *The Art of Computer Programming, Volume 1: Fundamental Algorithms*. 3rd ed. Addison-Wesley, 1997, pp. 435–441.
- [6] Z. Meng and S. Zhang. “Buddy Algorithm Optimization in Linux Memory Management.” In: *Applied Mechanics and Materials* 423–426 (2013), pp. 2746–2750.
- [7] J. L. Peterson and T. A. Norman. “Buddy Systems.” In: *Communications of the ACM* 20.6 (1977), pp. 421–431.
- [8] K. K. Shen and J. L. Peterson. “A Weighted Buddy Method for Dynamic Storage Allocation.” In: *Communications of the ACM* 17.10 (1974), pp. 558–563.
- [9] M. R. Q. Tandjung. *Applications of Divide and Conquer in the Buddy System for Efficient Memory Allocations*. IF2211 Strategi Algoritma, ITB. 2024.
- [10] P. R. Wilson et al. “Dynamic Storage Allocation: A Survey and Critical Review.” In: (1995).